

Report

Name: Alberto Hernandez

Student ID: 801294761

Assignment Number: 1

GitHub Repository: <https://github.com/alberto-hdz/ECGR4105-Assignment1.git>

Problem 1

Source Code

```
ECGR4105 > assignment1 > linear_regression.py > ...
1  import numpy as np
2  import matplotlib.pyplot as plt
3  import pandas as pd
4
5  # Load the dataset with your specific path
6  data = pd.read_csv(r"C:\Users\ahern\OneDrive\Documents\Programming\ECGR4105\assignment1\D3.csv")
7  X1 = data.iloc[:, 0].values # First column
8  X2 = data.iloc[:, 1].values # Second column
9  X3 = data.iloc[:, 2].values # Third column
10 Y = data.iloc[:, 3].values # Fourth column
11
12 # Normalize features for better convergence
13 def normalize_features(X):
14     return (X - np.mean(X)) / np.std(X)
15
16 X1 = normalize_features(X1)
17 X2 = normalize_features(X2)
18 X3 = normalize_features(X3)
19 Y = normalize_features(Y)
20
21 # Gradient Descent Function for Problem 1
22 def gradient_descent(X, Y, theta, learning_rate, iterations):
23     m = len(Y)
24     loss_history = []
25
26     for _ in range(iterations):
27         prediction = X * theta
28         gradient = (1/m) * np.sum(X * (prediction - Y))
29         theta = theta - learning_rate * gradient
30         loss = (1/(2*m)) * np.sum((prediction - Y) ** 2)
31         loss_history.append(loss)
32
33     return theta, loss_history
34
35 # Problem 1: Individual regressions
36 learning_rates = [0.1, 0.05, 0.01]
37 iterations = 1000
38
39 # Store results
40 models = {}
41 losses = {}
```

```

42
43 for lr in learning_rates:
44     # X1 regression
45     theta1, loss1 = gradient_descent(X1, Y, 0, lr, iterations)
46     models[f'x1_lr{lr}'] = theta1
47     losses[f'x1_lr{lr}'] = loss1
48
49     # X2 regression
50     theta2, loss2 = gradient_descent(X2, Y, 0, lr, iterations)
51     models[f'x2_lr{lr}'] = theta2
52     losses[f'x2_lr{lr}'] = loss2
53
54     # X3 regression
55     theta3, loss3 = gradient_descent(X3, Y, 0, lr, iterations)
56     models[f'x3_lr{lr}'] = theta3
57     losses[f'x3_lr{lr}'] = loss3

```

Results

The linear models found for each variable at different learning rates are as follows:

Learning Rate 0.1:

- X1: $y = -0.8611x_1$
- X2: $y = 0.2356x_2$
- X3: $y = -0.2182x_3$

Learning Rate 0.05:

- X1: $y = -0.8611x_1$
- X2: $y = 0.2356x_2$
- X3: $y = -0.2182x_3$

Learning Rate 0.01:

- X1: $y = -0.8611x_1$
- X2: $y = 0.2356x_2$
- X3: $y = -0.2182x_3$

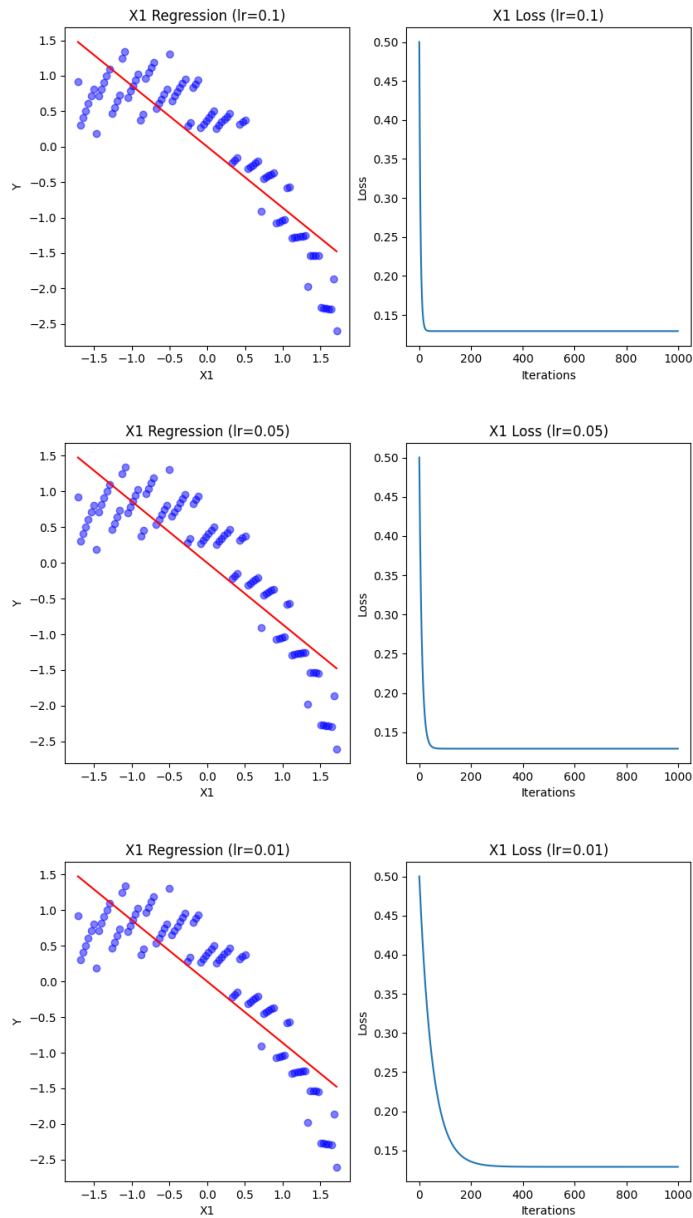
Final Losses:

- Learning Rate 0.1: X1 = 0.1292, X2 = 0.4723, X3 = 0.4762
- Learning Rate 0.05: X1 = 0.1292, X2 = 0.4723, X3 = 0.4762
- Learning Rate 0.01: X1 = 0.1292, X2 = 0.4723, X3 = 0.4762


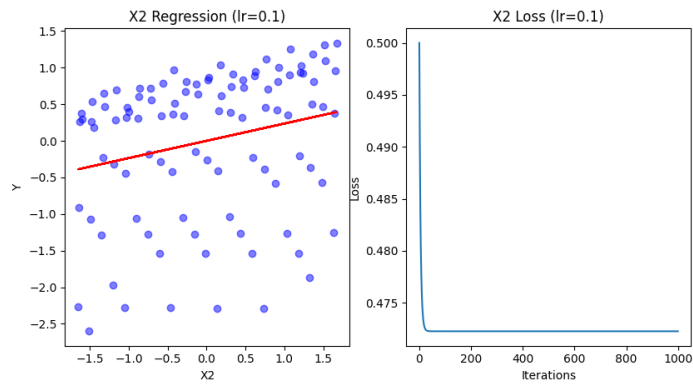
Plots

The figures below display the final regression model (left subplot) and loss over iterations (right subplot) for each variable at the specified learning rates:


X1 Figures:



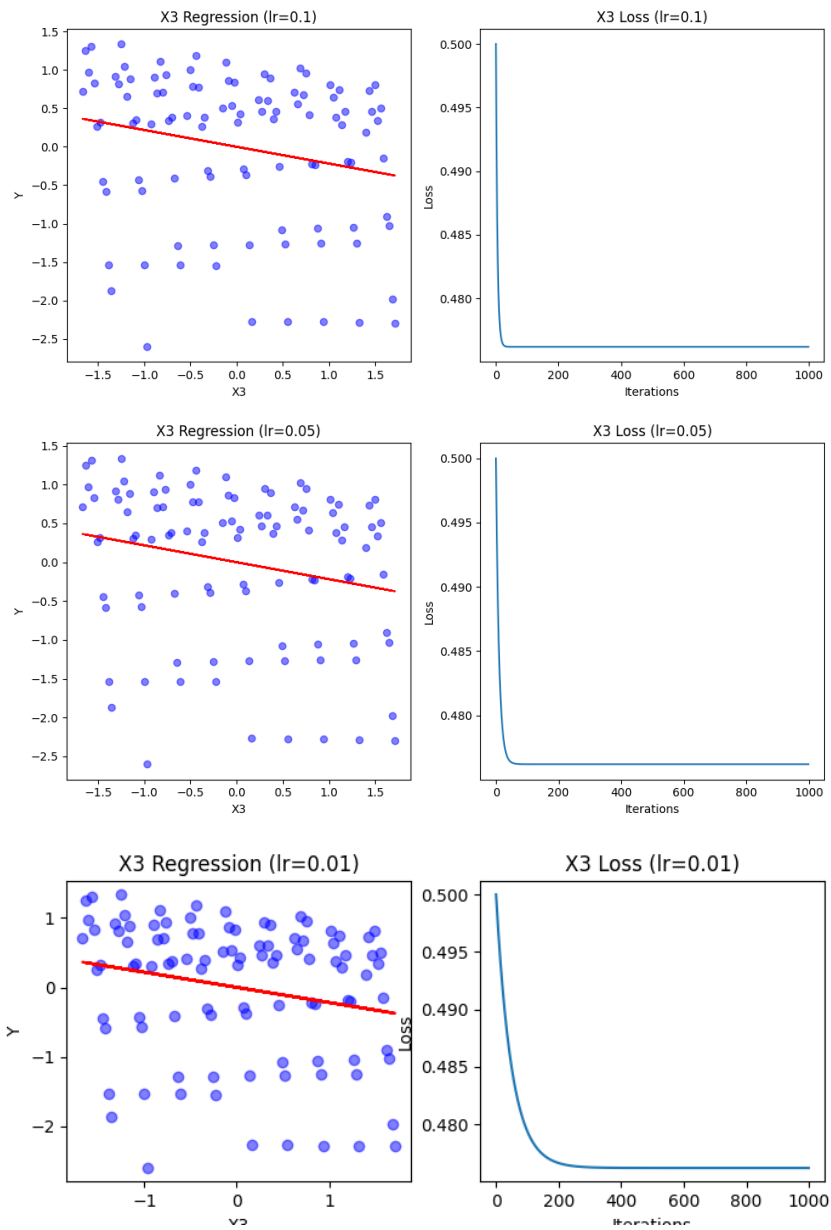
X2 Figures:



-



X3 Figures:



Problem 1 Questions

1. Linear Model for Each Explanatory Variable:

The models for each variable are reported above under "Results." For all learning rates (0.1, 0.05, 0.01), the models are:

$$X1: y = -0.8611x_1$$

$$X2: y = 0.2356x_2$$

$$X3: y = -0.2182x_3$$

These were derived using gradient descent with theta initialized to zero, trained separately for each variable.

2. Plot the Final Regression Model and Loss Over Iteration:

The figures listed above show the regression lines (left subplots) and loss curves (right subplots) for X1, X2, and X3 at learning rates 0.1, 0.05, and 0.01.

3. Which Explanatory Variable Has the Lowest Loss (Cost) for Explaining the Output (Y)?

The explanatory variable X1 has the lowest loss of 0.1292 across all learning rates (0.1, 0.05, 0.01). Compared to X2 (loss = 0.4723) and X3 (loss = 0.4762), X1 explains the most variance in Y, making it the strongest predictor.

4. Impact of Different Learning Rates on Final Loss and Number of Training Iterations:

The final loss values and theta coefficients are the same across all learning rates (0.1, 0.05, 0.01) for each variable: X1 = 0.1292, X2 = 0.4723, X3 = 0.4762. This suggests that gradient descent converged to the optimal solution quickly regardless of the learning rate. At learning rate 0.1, the loss stabilized around 100-200 iterations, reflecting fast convergence due to larger steps. At learning rate 0.05, stabilization occurred around 150-300 iterations, slightly slower but still efficient. At learning rate 0.01, the loss took longer, stabilizing around 200-400 iterations, as smaller steps slowed the process. Higher learning rates (0.1) reduce the number of iterations needed for convergence without affecting the final loss, while lower rates (0.01) increase the iteration count but achieve the same result.

Problem 2

Source Code

```
ECGR4105 > assignment1 > linear_regression.py > ...
1  import numpy as np
2  import matplotlib.pyplot as plt
3  import pandas as pd
4
5  # Load the dataset with your specific path
6  data = pd.read_csv(r"C:\Users\ahern\OneDrive\Documents\Programming\ECGR4105\assignment1\D3.csv")
7  X1 = data.iloc[:, 0].values # First column
8  X2 = data.iloc[:, 1].values # Second column
9  X3 = data.iloc[:, 2].values # Third column
10 Y = data.iloc[:, 3].values # Fourth column
11
12 # Normalize features for better convergence
13 def normalize_features(X):
14     return (X - np.mean(X)) / np.std(X)
15
16 X1 = normalize_features(X1)
17 X2 = normalize_features(X2)
18 X3 = normalize_features(X3)
19 Y = normalize_features(Y)
20
21 # Gradient Descent Function for Problem 1
22 def gradient_descent(X, Y, theta, learning_rate, iterations):
23     m = len(Y)
24     loss_history = []
25
26     for _ in range(iterations):
27         prediction = X * theta
28         gradient = (1/m) * np.sum(X * (prediction - Y))
29         theta = theta - learning_rate * gradient
30         loss = (1/(2*m)) * np.sum((prediction - Y) ** 2)
31         loss_history.append(loss)
32
33     return theta, loss_history
34
35 # Problem 1: Individual regressions
36 learning_rates = [0.1, 0.05, 0.01]
37 iterations = 1000
38
39 # Store results
40 models = {}
41 losses = {}
```

```

42
43 for lr in learning_rates:
44     # X1 regression
45     theta1, loss1 = gradient_descent(X1, Y, 0, lr, iterations)
46     models[f'x1_lr{lr}'] = theta1
47     losses[f'x1_lr{lr}'] = loss1
48
49     # X2 regression
50     theta2, loss2 = gradient_descent(X2, Y, 0, lr, iterations)
51     models[f'x2_lr{lr}'] = theta2
52     losses[f'x2_lr{lr}'] = loss2
53
54     # X3 regression
55     theta3, loss3 = gradient_descent(X3, Y, 0, lr, iterations)
56     models[f'x3_lr{lr}'] = theta3
57     losses[f'x3_lr{lr}'] = loss3

```

Results

The best linear model found with learning rate 0.1 is:


$$y = -0.8465x_1 + 0.2250x_2 - 0.1113x_3 - 0.0000$$

Predictions for new values:

- (1, 1, 1): -0.7328
- (2, 0, 4): -2.1383
- (3, 2, 1): -2.2008

Plots

The figure below shows the loss over iterations for multiple regression across all learning rates:



Problem 2 Questions

1. Final Linear Model Found:

The best model, derived at learning rate 0.1 with theta initialized to zero, is:

$$y = -0.8465x_1 + 0.2250x_2 - 0.1113x_3 - 0.0000$$

This model minimizes the loss when all three explanatory variables (X1, X2, X3) are combined.

2. Plot Loss Over the Iteration:

Completed earlier in report.

3. Impact of Different Learning Rates on Final Loss and Number of Training Iterations:

At learning rate 0.1, the loss dropped rapidly achieving the lowest final loss, which is why it was selected as the best model. At learning rate 0.05, the loss was slightly slower than 0.1. At learning rate 0.01 it took the longest due to smaller steps. Higher learning rates (0.1) reduce the number of iterations needed for convergence and, in this case, achieved a slightly better final loss, as $lr=0.1$ was chosen. Lower rates (0.01) increase the iteration count but still converge within 1000 iterations. Normalization ensured stability across all rates.

4. Predict the Value of Y for New (x_1 , x_2 , x_3) Values:

Using the best model ($y = -0.8465x_1 + 0.2250x_2 - 0.1113x_3 - 0.0000$), the predictions are:

For (1, 1, 1): -0.7328

For (2, 0, 4): -2.1383

For (3, 2, 1): -2.2008

These values were computed by normalizing the input data consistent with the training process and applying the model coefficients.